



# Locality-sensitive hashing without false negatives in $\ell_1$

Negev Shekel-Nosatzki, Chen-Yu Yang, Parthiban  
Loganathan

*Department of Computer Science  
Columbia University*

COMS 6998: Algorithmic Techniques for Massive Data

December 2015

- 1 Locality Sensitive Hashing Introduction
- 2 Our LSH construction
- 3 Conjecture
- 4 Future work

# Locality Sensitive Hashing

**Goal:** A data structure to distinguish between close and far points with near-linear space and sub-linear query time. Can be used in Nearest Neighbor Search for example.

## Definition (informal)

A *locality sensitive hash function*, *LSH* is a random hash function  $h : \mathbb{R}^d \rightarrow U$  ( $h$  drawn from a family  $\mathcal{H}$ ,  $U$  some finite set) such that

1.  $d(q, p) \leq r \implies \Pr[h(q) = h(p)] = P_1$  is “not-so-small”
2.  $d(q, p) > cr \implies \Pr[h(q) = h(p)] = P_2$  is “small”

Generally,  $P_1 < P_2$  and we use  $n^\rho$  hash tables where:

$$\rho = \frac{\log(1/P_1)}{\log(1/P_2)}.$$



# Locality Sensitive Hashing in Hamming

From class, LSH for Hamming Space  $\{0, 1\}^d$  with distance metric  $(x, y) = |\{x_i \neq y_i\}|$ .

The hash family,  $\mathcal{H} : \{g : \{0, 1\}^d \rightarrow \{0, 1\}^k\}$ , is defined as:

$$g(p) := (h_1(p), h_2(p), \dots, h_k(p)),$$

where

$$h_i(p) := p_j \text{ for random } j \leftarrow [d].$$

For NNS, we maintain  $L = n^\rho$  such Hamming-LSH tables  $g$ .



# LSH from Pagh'16

- ▶ It's **covering** - guarantees collision for  $\|x - y\| \leq r$ . ie. Probability of false negatives is 0.
- ▶ Far points **collide with low probability**.
- ▶ Chooses **correlated hash functions** instead of selecting them independently to cover all ways in which two points can be  $r$  apart.
- ▶ **Performance is almost as good as what we saw in class** where we had constant probability of false negatives. The algorithm is efficient if  $r = \Theta(\log n)$ .



# LSH from Pagh'16

- ▶ It's **covering** - guarantees collision for  $\|x - y\| \leq r$ . ie. Probability of false negatives is 0.
- ▶ Far points **collide with low probability**.
- ▶ Chooses **correlated hash functions** instead of selecting them independently to cover all ways in which two points can be  $r$  apart.
- ▶ **Performance is almost as good** as what we saw in class where we had constant probability of false negatives. The algorithm is efficient if  $r = \Theta(\log n)$ .



# LSH from Pagh'16

- ▶ It's **covering** - guarantees collision for  $\|x - y\| \leq r$ . ie. Probability of false negatives is 0.
- ▶ Far points **collide with low probability**.
- ▶ Chooses **correlated hash functions** instead of selecting them independently to cover all ways in which two points can be  $r$  apart.
- ▶ **Performance is almost as good** as what we saw in class where we had constant probability of false negatives. The algorithm is efficient if  $r = \Theta(\log n)$ .



# LSH from Pagh'16

- ▶ It's **covering** - guarantees collision for  $\|x - y\| \leq r$ . ie. Probability of false negatives is 0.
- ▶ Far points **collide with low probability**.
- ▶ Chooses **correlated hash functions** instead of selecting them independently to cover all ways in which two points can be  $r$  apart.
- ▶ **Performance is almost as good** as what we saw in class where we had constant probability of false negatives. The algorithm is efficient if  $r = \Theta(\log n)$ .





# Pagh's LSH Definition

[Pagh'16] Given  $S \subseteq \{0, 1\}^d$ ,  $c > 1$  and  $r \in \mathbf{N}$  there exists a data structure with parameter  $\rho$  such that:

- ▶ On query  $y \in \{0, 1\}^d$  the data structure is guaranteed to return  $x \in S$  with  $\|x - y\| < cr$  if there exists  $x' \in S$  with  $\|x' - y\| \leq r$ .
- ▶ The size is  $O(dn^{1+\rho})$  bits where  $n = |S|$ .
- ▶ The expected query time is  $O(n^\rho(1 + d/w))$ , where  $w$  is the word length.
- ▶ Parameter  $\rho$  is a function of  $n, r$  and  $c$  bounded by  $\rho \leq \min\left(\ln 4/c + \frac{\log r}{\log n}, 1/c + \frac{r}{\log n}\right)$ . If  $\log(n)/cr \in \mathbf{N}$ , then  $\rho = 1/c$ .



# Pagh LSH Construction

$\mathcal{H}_{\mathcal{A}} = \{x \mapsto x \wedge a \mid a \in \mathcal{A}\}$  where  $\mathcal{A} \subseteq \{0, 1\}^d$  is a set of bit masks. We go through all  $h \in \mathcal{H}_{\mathcal{A}}$  and check for collisions.

$\mathcal{A}(m) = \{a(v) \mid v \in \{0, 1\}^{r+1} \setminus \{\mathbf{0}\}\}$  where  $a(v)_i = \langle m(i), v \rangle \forall v \in \{0, 1\}^{r+1}$  and  $m : \{1, \dots, d\} \rightarrow \{0, 1\}^{r+1}$ .

Pick random function  $m$  uniformly.

Example: this procedure produces the following covering set when  $m(i)$  is the binary representation of  $i$ :

$$\begin{pmatrix} 1 & 0 & 1 & 0 & 1 & 0 & 1 \\ 0 & 1 & 1 & 0 & 0 & 1 & 1 \\ 1 & 1 & 0 & 0 & 1 & 1 & 0 \\ 0 & 0 & 0 & 1 & 1 & 1 & 1 \\ 1 & 0 & 1 & 1 & 0 & 1 & 0 \\ 0 & 1 & 1 & 1 & 1 & 0 & 0 \\ 1 & 1 & 0 & 1 & 0 & 0 & 1 \end{pmatrix}$$

1 samples the coordinate. More efficient than having  $\binom{d}{r}$  possible values for  $a$ .



# $\ell_1$ LSH w/o false negatives

Goal: Create an efficient Hashing that always collide close points and rarely collide far points

Plan:   ▶ Embed  $\ell_1$  into Hamming  
          ▶ Hash Hamming

Embed: Assuming  $x \in [M]^d$ , we can embed  $x$  to  $\{0, 1\}^{Md}$  using unary embedding (as we saw in PS4) and hash that.

Eff.:   ▶  $T = \mathcal{O}(M \cdot d \cdot n^\rho)$   
          ▶  $S = \mathcal{O}(M \cdot d \cdot n^{1+\rho})$

Can we do better?



# $\ell_1$ LSH w/o false negatives

**Goal:** Create an efficient Hashing that always collide close points and rarely collide far points

**Plan:**

- ▶ Embed  $\ell_1$  into Hamming
- ▶ Hash Hamming

**Embed:** Assuming  $x \in [M]^d$ , we can embed  $x$  to  $\{0, 1\}^{Md}$  using unary embedding (as we saw in PS4) and hash that.

**Eff.:**

- ▶  $T = \mathcal{O}(M \cdot d \cdot n^\rho)$
- ▶  $S = \mathcal{O}(M \cdot d \cdot n^{1+\rho})$

Can we do better?



# $\ell_1$ LSH w/o false negatives

**Goal:** Create an efficient Hashing that always collide close points and rarely collide far points

**Plan:**

- ▶ Embed  $\ell_1$  into Hamming
- ▶ Hash Hamming

**Embed:** Assuming  $x \in [M]^d$ , we can embed  $x$  to  $\{0, 1\}^{Md}$  using unary embedding (as we saw in PS4) and hash that.

**Eff.:**

- ▶  $T = \mathcal{O}(M \cdot d \cdot n^\rho)$
- ▶  $S = \mathcal{O}(M \cdot d \cdot n^{1+\rho})$

Can we do better?



# $\ell_1$ LSH w/o false negatives

**Goal:** Create an efficient Hashing that always collide close points and rarely collide far points

**Plan:**

- ▶ Embed  $\ell_1$  into Hamming
- ▶ Hash Hamming

**Embed:** Assuming  $x \in [M]^d$ , we can embed  $x$  to  $\{0,1\}^{Md}$  using unary embedding (as we saw in PS4) and hash that.

**Eff.:**

- ▶  $T = \mathcal{O}(M \cdot d \cdot n^\rho)$
- ▶  $S = \mathcal{O}(M \cdot d \cdot n^{1+\rho})$

Can we do better?



# $\ell_1$ LSH w/o false negatives

**Goal:** Create an efficient Hashing that always collide close points and rarely collide far points

**Plan:**

- ▶ Embed  $\ell_1$  into Hamming
- ▶ Hash Hamming

**Embed:** Assuming  $x \in [M]^d$ , we can embed  $x$  to  $\{0,1\}^{Md}$  using unary embedding (as we saw in PS4) and hash that.

**Eff.:**

- ▶  $T = \mathcal{O}(M \cdot d \cdot n^\rho)$
- ▶  $S = \mathcal{O}(M \cdot d \cdot n^{1+\rho})$

Can we do better?



# $\ell_1$ LSH w/o false negatives

**Goal:** Create an efficient Hashing that always collide close points and rarely collide far points

**Plan:**

- ▶ Embed  $\ell_1$  into Hamming
- ▶ Hash Hamming

**Embed:** Assuming  $x \in [M]^d$ , we can embed  $x$  to  $\{0,1\}^{Md}$  using unary embedding (as we saw in PS4) and hash that.

**Eff.:**

- ▶  $T = \mathcal{O}(M \cdot d \cdot n^\rho)$
- ▶  $S = \mathcal{O}(M \cdot d \cdot n^{1+\rho})$

Can we do better?





# $\ell_1$ Embedding Setup

- ▶ Can we use Sampling ?
  - ▶ Random Sampling will miss far points too often
  - ▶ Grid shifting can cut close points - exponential (in  $r$  or  $d$ ) cost to avoid that prob.
- ▶ **Our Solution - Modulo Embedding**
  - ▶ Keep close points close
  - ▶ Contract distance of far points - but keep them far enough ( $> cr$ ) with constant probability
  - ▶ Amplify probability by iterating and concatenating



# $\ell_1$ Embedding Setup

- ▶ Can we use Sampling ?
  - ▶ Random Sampling will miss far points too often
  - ▶ Grid shifting can cut close points - exponential (in  $r$  or  $d$ ) cost to avoid that prob.
- ▶ **Our Solution - Modulo Embedding**
  - ▶ Keep close points close
  - ▶ Contract distance of far points - but keep them far enough ( $> cr$ ) with constant probability
  - ▶ Amplify probability by iterating and concatenating



# $\ell_1$ Embedding Setup

- ▶ Can we use Sampling ?
  - ▶ Random Sampling will miss far points too often
  - ▶ Grid shifting can cut close points - exponential (in  $r$  or  $d$ ) cost to avoid that prob.
- ▶ **Our Solution - Modulo Embedding**
  - ▶ Keep close points close
  - ▶ Contract distance of far points - but keep them far enough ( $> cr$ ) with constant probability
  - ▶ Amplify probability by iterating and concatenating



# Modulo Mapping

- Define  $f_\alpha : \mathbb{N} \rightarrow \{0,1\}^\alpha$  to be the following mapping with parameter  $\alpha$ :

$$f_\alpha(x) = \begin{cases} 0^{\alpha-\beta} 1^\beta, & \text{if } \beta < \alpha \\ 1^{2\alpha-\beta} 0^{\beta-\alpha}, & \text{otherwise} \end{cases}$$

Where  $\beta = x \pmod{2\alpha}$ .

- Example with  $\alpha = 5$ :

$\beta = 0 \rightarrow 00000$      $\beta = 5 \rightarrow 11111$

$\beta = 1 \rightarrow 00001$      $\beta = 6 \rightarrow 11110$

$\beta = 2 \rightarrow 00011$      $\beta = 7 \rightarrow 11100$

$\beta = 3 \rightarrow 00111$      $\beta = 8 \rightarrow 11000$

$\beta = 4 \rightarrow 01111$      $\beta = 9 \rightarrow 10000$

- $f_\alpha$  preserves modulo distance !

$$\begin{aligned} \|f_\alpha(x) - f_\alpha(y)\|_{Ham} &= \|f_\alpha(|x - y|)\|_1 \\ &= \min(x - y \pmod{2\alpha}, y - x \pmod{2\alpha}) \end{aligned}$$



# Modulo Mapping

- Define  $f_\alpha : \mathbb{N} \rightarrow \{0,1\}^\alpha$  to be the following mapping with parameter  $\alpha$ :

$$f_\alpha(x) = \begin{cases} 0^{\alpha-\beta} 1^\beta, & \text{if } \beta < \alpha \\ 1^{2\alpha-\beta} 0^{\beta-\alpha}, & \text{otherwise} \end{cases}$$

Where  $\beta = x \pmod{2\alpha}$ .

- Example with  $\alpha = 5$ :

$$\beta = 0 \rightarrow 00000 \quad \beta = 5 \rightarrow 11111$$

$$\beta = 1 \rightarrow 00001 \quad \beta = 6 \rightarrow 11110$$

$$\beta = 2 \rightarrow 00011 \quad \beta = 7 \rightarrow 11100$$

$$\beta = 3 \rightarrow 00111 \quad \beta = 8 \rightarrow 11000$$

$$\beta = 4 \rightarrow 01111 \quad \beta = 9 \rightarrow 10000$$

- $f_\alpha$  **preserves modulo distance !**

$$\begin{aligned} \|f_\alpha(x) - f_\alpha(y)\|_{Ham} &= \|f_\alpha(|x - y|)\|_1 \\ &= \min(x - y \pmod{2\alpha}, y - x \pmod{2\alpha}) \end{aligned}$$



## $\ell_1 \rightarrow$ Hamming Modulo Embedding - Construction

- ▶ Define  $T : [M]^d \rightarrow \{0, 1\}^*$  as follows:
  - ▶ Randomly pick  $k = \log(n)$  integers  $\alpha_1.. \alpha_k$  from  $[A, 2A]$ , where  $A = \max(5cr, \log(M))$
  - ▶ Run  $f_{\alpha_j}(x_i)$  for each  $i \in [d], j \in [k]$
  - ▶ Concatenate all strings
- ▶ Given input  $x \in [M]^d$ , we will hash  $T(x)$  using Pagh with  $r' = k \cdot r$



# $\ell_1 \rightarrow$ Hamming Modulo Embedding - Analysis

- ▶ Final Dimension:  $\mathcal{O}(d \cdot \log(n) \cdot (r + \log(M)))$  (but also now we are in bit-space, so representation and computation are  $\log(M)$  faster).
  - ▶  $T = \mathcal{O}((r + \log(M)) \cdot \log(n) \cdot d \cdot n^\rho)$
  - ▶  $S = \mathcal{O}((r + \log(M)) \cdot \log(n) \cdot d \cdot n^{1+\rho})$
- ▶ When  $r = \Theta(\log(n))$  (this is where Pagh focus on and when his algo is optimal and  $\rho = 1/c$ ) - our expected query time is  $T = \mathcal{O}(\text{polylog}(n, M) \cdot d \cdot n^{1/c})$  which is much faster than any existing deterministic  $\ell_1$  NNS algorithm with near-linear space.



# Conjecture behind Construction

- ▶ Conjecture: The modulo function keeps close points close all the time, and any far point will stay far on average. Specifically:

$$\forall x > 1, \forall A \geq \log(x), \alpha \sim \mathcal{U}[A, 2A] :$$

$$\mathbb{E}_{\alpha}[\min(x \bmod 2\alpha, -x \bmod 2\alpha)] \geq \min(x, \epsilon\alpha)$$

For some constant  $\epsilon$

- ▶ As long as we sample  $\alpha_i \geq 2cr/\epsilon$ , then expectation for far coordinates is larger than  $2cr$ .
- ▶ Variance is bounded by range
- ▶ After  $\log(n)$  times  $O(1)$  far points will become close
- ▶ Cannot be 100% correct for any  $A$  (think about  $x = LCM(\alpha_1 \dots \alpha_k) + 1$ )
  - ▶ It's okay since  $LCM(\alpha_1 \dots \alpha_k) > M$





# Test of conjecture

- ▶ We run the following experiments
- ▶ **Completeness Test**
  - ▶ Test all  $[2^{26}]$ ,  $K = 20$ ,  $cr \in [4, 50]$ .
  - ▶ No far points becomes close in  $T = 10$  rounds for each  $cr$ .
- ▶ **Big Numbers Test**
  - ▶ Each round randomly sample  $2^{20}$  numbers in  $[2^{60}]$ ,  $K = 20$ ,  $cr \in \{4, 10, 50, 100\}$ .
  - ▶ No far points becomes close in  $T = 100$  rounds for each  $cr$ .



# What if $r$ is big?

- ▶ If we need to hash with big  $r$ , we can replace  $r$  factor with  $d$  and pay constant growth in exponent and denominator (trade-off)
- ▶ Divide each coordinate by  $\frac{r\Delta}{d}$  (for some constant  $\Delta < c - 1$ ) and round them to the nearest integer.
- ▶ far point might contract to  $\frac{cr}{1+\Delta}$ , so we hash (using Pagh) with  $c' = \frac{c}{1+\Delta}$ .
- ▶ Final dimension becomes  $\mathcal{O}(\log(n)(d^2 + d \log(M))/\Delta)$  and the time becomes  $\approx n^{\rho+\Delta}$ .



# Acknowledgements

- ▶ Alex for helping us with our project direction.
- ▶ Pagh, Rasmus. "Locality-sensitive hashing without false negatives." arXiv preprint arXiv:1507.03225 (2015).



End

Thank you

Questions?

